

**AWS CLOUD COMPUTING COHORT 2**

**CAPSTONE PROJECT 2**

**SERVERLESS INFRASTRUCTURE-AS-CODE (IAC)  
SOLUTION FOR AUTOMATED LANGUAGE  
TRANSLATION ON AWS**

**EUGENE QUAYE TETTEH**

**AZUBI AFRICA  
5<sup>TH</sup> SEPTEMBER, 2025**

## Table of Contents

### Table of Contents

Project Overview .....	4
Objectives.....	4
System Architecture .....	4
Architecture Diagram Description .....	5
Cloud Services Selection and Justification .....	5
Amazon S3.....	5
AWS Lambda .....	6
AWS Translate .....	6
AWS IAM .....	6
Terraform .....	6
Prerequisites .....	6
Infrastructure Implementation with Terraform .....	7
S3 Buckets .....	7
IAM Policies & Roles.....	7
Lambda Function.....	7
S3 Bucket Notification.....	7
Providers and Variables (variables.tf) .....	8
S3 Buckets with Lifecycle Policies (main.tf).....	9
Lambda IAM Role and Policy (iam.tf).....	10
Lambda Function (lambda.tf).....	11
S3 Event Trigger (event_notification.tf) .....	11
Lambda Function Development & Packaging .....	12
Python Lambda Function (lambda_function.py) .....	12
Packaging Instructions: .....	13
End-to-End AWS Deployment and Teardown Using Terraform .....	13
AWS CLI Configuration .....	13
Installing Boto3 .....	14

Packaging Lambda Function.....	14
Terraform Initialized .....	14
Terraform Planning.....	15
Terraform Apply .....	16
Resource Creation Confirmation.....	16
Testing & Validation .....	25
Sample Input .....	25
Expected Output .....	25
Challenges .....	27
Lessons learned.....	27
Cost management.....	28
Conclusion.....	28
References.....	29

# Project Overview

This project delivers a fully serverless, Infrastructure-as-Code (IaC) solution for automated language translation on AWS, using CloudFormation or Terraform to provision resources. It leverages AWS Translate for NLP processing and Amazon S3 for object storage. Translation requests in JSON format are processed using a Python script (Boto3), which stores original and translated data in separate S3 buckets. Automation is achieved with AWS Lambda, making this solution scalable, cost-effective, and easily extensible, while remaining within AWS Free Tier limits.

## Objectives

- Automate end-to-end language translation using scalable, serverless AWS services, requiring no manual infrastructure management.
- Enable Continuous Deployment: All cloud resources are provisioned, updated, and decommissioned automatically using Terraform.
- Guarantee Security: Employ robust IAM roles and policies adhering to the principle of least privilege.
- Maintain Cost-Effectiveness: Operate entirely within AWS Free Tier, using lifecycle management to minimize S3 costs.
- Create Reproducible Results: All configurations, policies, and code are tracked and can be re-deployed or destroyed in any AWS account.

## System Architecture

A user uploads a set of texts (as a JSON file) for translation. Upon upload, the system immediately processes the file:

Amazon S3 receives the JSON file (source texts plus source and target language codes).

An S3 Event Notification triggers the Lambda function.

The Lambda function reads the file, translates each text using AWS Translate, and writes a new JSON with translations to an S3 response bucket.

The user retrieves the translated file from S3.

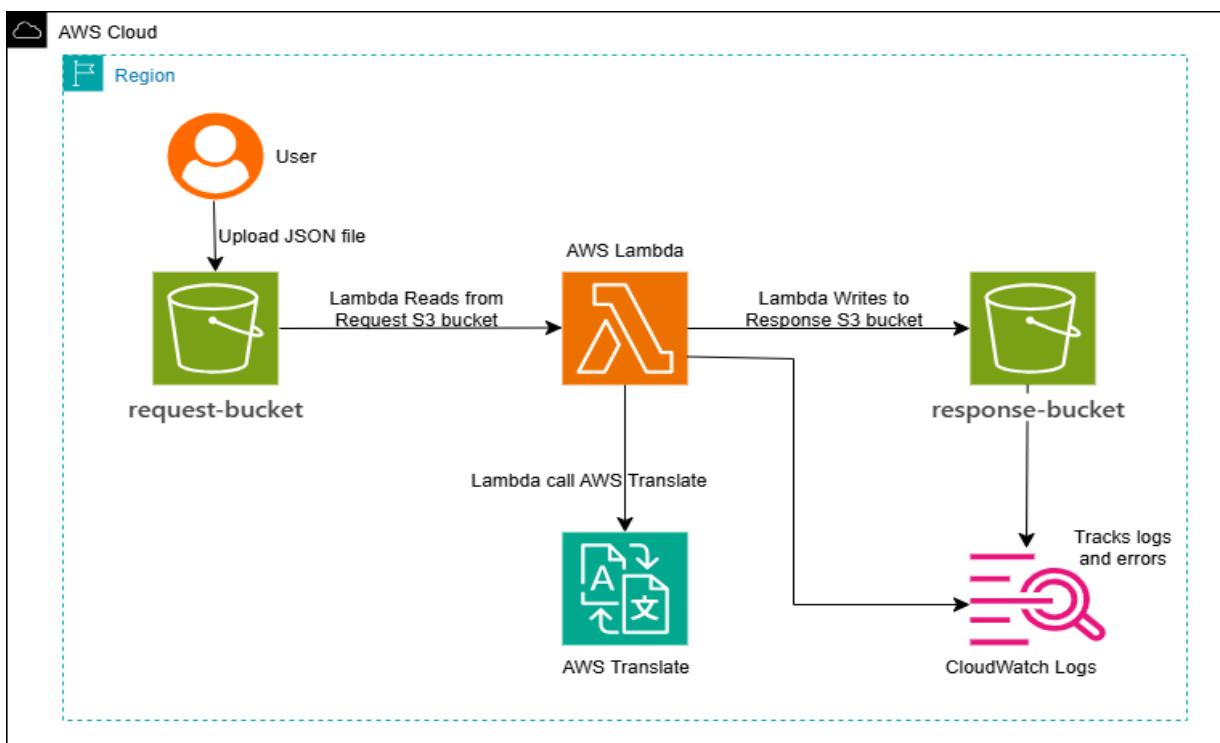
## Architecture Diagram Description:

S3 (request-bucket) ← [Upload JSON]

S3 Event → Lambda (translate-json-files)

Lambda → AWS Translate API

Lambda → S3 (response-bucket) [Upload translated JSON]



## Cloud Services Selection and Justification

### Amazon S3

Use: Stores input files (to be translated) and output files (translated text).

Reason: Simple object storage, supports event notifications, Free Tier eligible.

## AWS Lambda

Use: Serverless compute to handle all translation logic on-demand.

Reason: No server management; scales instantly with usage; 1M free executions/month.

## AWS Translate

Use: Performs the actual language translation.

Reason: Pre-trained, neural machine translation, Free Tier up to 2M characters/month.

## AWS IAM

Use: Restricts Lambda and S3 operations to the minimum set required.

Reason: Prevents security breaches, follows security best practices.

## Terraform

Use: Defines entire system state as code — provisioning, updating, and destroying cloud resources as a unified workflow.

Reason: Reproducibility, version control, and automation, which are industry best practice for cloud projects.

## Prerequisites

- AWS Account (Free Tier eligible)
- AWS CLI and credentials configured (aws configure)
- Python 3.7+ installed locally
- Terraform v1.x installed
- Boto3 Python package
- Basic knowledge of AWS IAM, Lambda, S3

# Infrastructure Implementation with Terraform

## S3 Buckets

Two buckets: request-bucket (input) and response-bucket (output).

Both have lifecycle rules to auto-delete objects after 7 days, ensuring no unexpected charges.

## IAM Policies & Roles

One role for Lambda, with a custom inline policy granting only S3 and Translate API access.

S3 buckets are locked down with only this role having access.

## Lambda Function

Deployed using Terraform with all environment variables configured.

Source code packaged as a .zip and referenced by the Lambda resource.

## S3 Bucket Notification

Provisions an S3 event on object creation to invoke Lambda.

Ensures Lambda is only triggered when new JSON files are uploaded.

**Below are the key Terraform resources:**

Providers and Variables (variables.tf)

```
variables.tf > variable "response_bucket"
1 ✓ provider "aws" {
2   |   region = var.aws_region
3   }
4
5 ✓ variable "aws_region" {
6   |   default = "us-east-1"
7   }
8 ✓ variable "request_bucket" {
9   |   default = "eugene-unique-request-bucket"
10  }
11 ✓ variable "response_bucket" [
12   |   default = "eugene-unique-response-bucket"
13  ]
```

## S3 Buckets with Lifecycle Policies (main.tf)

```
main.tf > resource "aws_s3_bucket" "request" > tags > Name
1  resource "aws_s3_bucket" "request" {
2    bucket      = var.request_bucket
3    force_destroy = true
4
5    tags = [
6      Environment = "Dev"
7      Name        = "Request Bucket"
8    ]
9  }
10
11 resource "aws_s3_bucket_lifecycle_configuration" "request.lifecycle" {
12   bucket = aws_s3_bucket.request.id
13   rule {
14     id      = "auto-delete"
15     status = "Enabled"
16
17     filter {
18       prefix = ""
19     }
20
21     expiration {
22       days = 7
23     }
24   }
25 }
26
27 resource "aws_s3_bucket" "response" {
28   bucket      = var.response_bucket
29   force_destroy = true
30
31   tags = [
32     Environment = "Dev"
33     Name        = "Response Bucket"
34   ]
35 }
36
37 resource "aws_s3_bucket_lifecycle_configuration" "response.lifecycle" {
38   bucket = aws_s3_bucket.response.id
39
40   rule {
41     id      = "auto-delete"
42     status = "Enabled"
43
44     filter {
45       prefix = ""
46     }
47
48     expiration {
```

## Lambda IAM Role and Policy (iam.tf)

```
iam.tf > ↵ resource "aws_iam_policy" "lambda_policy" > [e] policy
  1  data "aws_iam_policy_document" "lambda_assume_role" {
  2    statement {
  3      actions = ["sts:AssumeRole"]
  4      principals {
  5        type     = "Service"
  6        identifiers = ["lambda.amazonaws.com"]
  7      }
  8    }
  9  }
10
11 resource "aws_iam_role" "lambda_exec" {
12   name          = "lambda_translate_exec"
13   assume_role_policy = data.aws_iam_policy_document.lambda_assume_role.json
14 }
15
16 resource "aws_iam_policy" "lambda_policy" {
17   name      = "lambda_translate_policy"
18   description = "Allows Lambda to use Translate and access specific S3 buckets"
19   policy    = jsonencode({
20     Version = "2012-10-17",
21     Statement = [
22       {
23         Effect = "Allow",
24         Action = [
25           "translate:TranslateText"
26         ],
27         Resource = "*"
28       },
29       {
30         Effect = "Allow",
31         Action = [
32           "s3:GetObject",
33           "s3:PutObject",
34           "s3>ListBucket"
35         ],
36         Resource = [
37           aws_s3_bucket.request.arn,
38           "${aws_s3_bucket.request.arn}/*",
39           aws_s3_bucket.response.arn,
40           "${aws_s3_bucket.response.arn}/*"
41         ]
42       },
43       {
44         Effect = "Allow",
45         Action = [
46           "logs>CreateLogGroup",
47           "logs>CreateLogStream",
48           "logs>PutLogEvents"
49         ]
50       }
51     ]
52   }
53 }
```

## Lambda Function (lambda.tf)

```
λ lambda.tf > 📂 resource "aws_lambda_function" "translate_function" > 📂 environment
1   resource "aws_lambda_function" "translate_function" {
2     filename        = "lambda_function_payload.zip"
3     function_name  = "translate-json-files"
4     role           = aws_iam_role.lambda_exec.arn
5     handler        = "lambda_function.lambda_handler"
6     runtime         = "python3.9"
7     timeout         = 60
8     source_code_hash = filebase64sha256("lambda_function_payload.zip")
9     environment {
10       variables = {
11         RESPONSE_BUCKET = var.response_bucket
12       }
13     }
14 }
```

## S3 Event Trigger (event\_notification.tf)

```
λ event_noification.tf > 📂 resource "aws_s3_bucket_notification" "request_trigger"
1   resource "aws_s3_bucket_notification" "request_trigger" {
2     bucket = aws_s3_bucket.request.id
3
4     lambda_function {
5       lambda_function_arn = aws_lambda_function.translate_function.arn
6       events            = ["s3:ObjectCreated:*"]
7       filter_prefix      = "requests/"
8       filter_suffix      = ".json"
9     }
10 }
11
12 resource "aws_lambda_permission" "allow_s3" {
13   statement_id  = "AllowExecutionFromS3"
14   action        = "lambda:InvokeFunction"
15   function_name = aws_lambda_function.translate_function.function_name
16   principal    = "s3.amazonaws.com"
17   source_arn    = aws_s3_bucket.request.arn
18 }
```

## Lambda Function Development & Packaging

### Python Lambda Function (lambda\_function.py)

```
❶ lambda_function.py
1  import json
2  import boto3
3  import os
4
5  s3 = boto3.client('s3')
6  translate = boto3.client('translate')
7  response_bucket = os.environ.get('RESPONSE_BUCKET')
8
9  def lambda_handler(event, context):
10      # Extract bucket & key info from trigger
11      bucket = event['Records'][0]['s3']['bucket']['name']
12      key = event['Records'][0]['s3']['object']['key']
13      # Download and process input json
14      obj = s3.get_object(Bucket=bucket, Key=key)
15      data = json.loads(obj['Body'].read().decode('utf-8'))
16
17      source_lang = data.get('source_lang')
18      target_lang = data.get('target_lang')
19      texts      = data.get('texts', [])
20      translated = []
21      for line in texts:
22          result = translate.translate_text(
23              Text=line,
24              SourceLanguageCode=source_lang,
25              TargetLanguageCode=target_lang
26          )
27          translated.append(result['TranslatedText'])
28      # Prepare output JSON
29      out_json = {
30          "input_language": source_lang,
31          "output_language": target_lang,
32          "original": texts,
33          "translated": translated
34      }
35      # Write result to response bucket
36      out_key = key.replace('requests/', '').replace('.json', '_translated.json')
37      s3.put_object(
38          Bucket=response_bucket,
39          Key=f"responses/{out_key}",
40          Body=json.dumps(out_json).encode('utf-8')
41      )
42      return {"status": "done", "output": out_key}
```

You must upload the packaged *lambda\_function\_payload.zip* and reference it locally.

## Packaging Instructions:

1. Place the Python script as *lambda\_function.py* in a folder
2. cd to the folder and run: “**Compress-Archive -Path lambda\_function.py -DestinationPath "C:\Users\quaye\Desktop\Capstone Project 2\lambda\_function\_payload.zip"**
3. Make sure you update the Lambda resource in Terraform with the correct zip path (*filename*).

## End-to-End AWS Deployment and Teardown Using Terraform

### AWS CLI Configuration

Initiated the AWS CLI setup using the command: *aws configure*

Provided the required credentials:

- Access Key ID
- Secret Access Key
- Default Region Name
- Output Format

```
PS C:\Users\quaye\Desktop\Capstone Project 2> aws configure
AWS Access Key ID [*****QIP]: *****
AWS Secret Access Key [*****zpNm]: *****
Default region name [us-east-1]:
Default output format [json]:
```

## Installing Boto3

Installed the AWS SDK for Python (Boto3) to enable programmatic interaction with AWS services:

```
PS C:\Users\quaye\Desktop\Capstone Project 2> pip install boto3
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: boto3 in c:\users\quaye\appdata\roaming\python\python313\site-packages (1.40.21)
Requirement already satisfied: botocore<1.41.0,>=1.40.21 in c:\users\quaye\appdata\roaming\python\python313\site-packages (from boto3) (1.40.21)
Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in c:\users\quaye\appdata\roaming\python\python313\site-packages (from boto3) (1.0.1)
Requirement already satisfied: s3transfer<0.14.0,>=0.13.0 in c:\users\quaye\appdata\roaming\python\python313\site-packages (from boto3) (0.13.1)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in c:\users\quaye\appdata\roaming\python\python313\site-packages (from botocore<1.41.0,>=1.40.21->boto3) (2.9.0.post0)
Requirement already satisfied: urllib3!=2.2.0,<3,>=1.25.4 in c:\users\quaye\appdata\roaming\python\python313\site-packages (from botocore<1.41.0,>=1.40.21->boto3) (2.5.0)
Requirement already satisfied: six>=1.5 in c:\users\quaye\appdata\roaming\python\python313\site-packages (from python-dateutil<3.0.0,>=2.1->botocore<1.41.0,>=1.40.21->boto3) (1.17.0)
```

## Packaging Lambda Function

Compressed the Lambda function script into a ZIP archive for deployment:

```
PS C:\Users\quaye\Desktop\Capstone Project 2> Compress-Archive -Path lambda_function.py -DestinationPath "C:\Users\quaye\Desktop\Capstone Project 2\lambda_function_payload.zip"
»
```

## Terraform Initialized

The terraform init command was successfully executed, setting up the working directory and installing required provider plugins.

```
PS C:\Users\quaye\Desktop\Capstone Project 2> terraform init
Initializing the backend...
Initializing provider plugins...
- Finding latest version of hashicorp/aws...
- Installing hashicorp/aws v6.11.0...
- Installed hashicorp/aws v6.11.0 (signed by HashiCorp)
Terraform has created a lock file .terraform.lock.hcl to record the provider selections it made above. Include this file in your version control repository so that Terraform can guarantee to make the same selections by default when you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see any changes that are required for your infrastructure. All Terraform commands should now work.

If you ever set or change modules or backend configuration for Terraform, rerun this command to reinitialize your working directory. If you forget, other commands will detect it and remind you to do so if necessary.
```

## Terraform Planning

Created an execution plan to preview the changes Terraform would make:

```
PS C:\Users\quaye\Desktop\Capstone Project 2> terraform plan
data.aws_iam_policy_document.lambda_assume_role: Reading...
data.aws_iam_policy_document.lambda_assume_role: Read complete after 0s [id=2690255455]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_iam_policy.lambda_policy will be created
+ resource "aws_iam_policy" "lambda_policy" {
    + arn          = (known after apply)
    + attachment_count = (known after apply)
    + description   = "Allows Lambda to use Translate and access specific S3 buckets"
    + id           = (known after apply)
    + name         = "lambda_translate_policy"
    + name_prefix  = (known after apply)
    + path          = "/"
    + policy        = (known after apply)
    + policy_id    = (known after apply)
    + tags_all     = (known after apply)
}

# aws_iam_role.lambda_exec will be created
+ resource "aws_iam_role" "lambda_exec" {
    + arn          = (known after apply)
    + assume_role_policy = jsonencode(
        {
            + Statement = [
                +
                    {
                        + Action    = "sts:AssumeRole"
                        + Effect   = "Allow"
                        + Principal = {
                            + Service = "lambda.amazonaws.com"
                        }
                    }
                ],
            + Version   = "2012-10-17"
        }
    )
    + create_date      = (known after apply)
    + force_detach_policies = false
    + id              = (known after apply)
    + managed_policy_arns = (known after apply)
    + max_session_duration = 3600
    + name            = "lambda translate exec"
}
```

## Terraform Apply

Applied the Terraform configuration to provision AWS resources. Confirmed the action by entering yes when prompted:

```
PS C:\Users\quaye\Desktop\Capstone Project 2> terraform apply
data.aws_iam_policy_document.lambda_assume_role: Reading...
data.aws_iam_policy_document.lambda_assume_role: Read complete after 0s [id=2690255455]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_iam_policy.lambda_policy will be created
+ resource "aws_iam_policy" "lambda_policy" {
    + arn          = (known after apply)
    + attachment_count = (known after apply)
    + description      = "Allows Lambda to use Translate and access specific S3 buckets"
    + id             = (known after apply)
    + name           = "lambda_translate_policy"
    + name_prefix     = (known after apply)
    + path            = "/"
    + policy          = (known after apply)
    + policy_id       = (known after apply)
    + tags_all        = (known after apply)
}

# aws_iam_role.lambda_exec will be created
+ resource "aws_iam_role" "lambda_exec" {
    + arn          = (known after apply)
    + assume_role_policy = jsonencode(
        {
            + Statement = [
                + {
                    + Action      = "sts:AssumeRole"
                    + Effect      = "Allow"
                    + Principal   = {
                        + Service = "lambda.amazonaws.com"
                    }
                },
            ],
        }
    )
}
```

## Resource Creation Confirmation

AWS resources were successfully created as defined in the Terraform configuration.

```
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

Enter a value: yes

aws_iam_role.lambda_exec: Creating...
aws_s3_bucket.request: Creating...
aws_s3_bucket.response: Creating...
aws_iam_role.lambda_exec: Creation complete after 2s [id=lambda_translate_exec]
aws_lambda_function.translate_function: Creating...
aws_s3_bucket.request: Creation complete after 7s [id=eugene-unique-response-bucket]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Creating...
aws_s3_bucket.request: Creation complete after 7s [id=eugene-unique-request-bucket]
aws_iam_policy.lambda_policy: Creating...
aws_s3_bucket.lifecycle_configuration.request.lifecycle: Creating...
aws_iam_policy.lambda_policy: Creation complete after 1s [id=arn:aws:iam::611837360746:policy/lambda_translate_policy]
aws_iam_role_policy_attachment.lambda_policy_attachment: Creating...
aws_iam_role_policy_attachment.lambda_policy_attachment: Creation complete after 1s [id=lambda_translate_exec/arn:aws:iam::611837360746:policy/lambda_translate_policy]
aws_lambda_function.translate_function: Still creating... [00m00s elapsed]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Still creating... [00m00s elapsed]
aws_s3_bucket.lifecycle_configuration.request.lifecycle: Still creating... [00m00s elapsed]
aws_lambda_function.translate_function: Creation complete after 20s [id=translate-json-files]
aws_lambda_permission.allow_s3: Creating...
aws_s3_bucket.notification.request_trigger: Creating...
aws_lambda_permission.allow_s3: Creation complete after 1s [id-AllowExecutionFromS3]
aws_s3_bucket.notification.request_trigger: Creation complete after 1s [id=eugene-unique-request-bucket]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Still creating... [00m00s elapsed]
aws_s3_bucket.lifecycle_configuration.request.lifecycle: Still creating... [00m20s elapsed]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Still creating... [00m30s elapsed]
aws_s3_bucket.lifecycle_configuration.request.lifecycle: Still creating... [00m30s elapsed]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Still creating... [00m40s elapsed]
aws_s3_bucket.lifecycle_configuration.request.lifecycle: Still creating... [00m40s elapsed]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Still creating... [00m50s elapsed]
aws_s3_bucket.lifecycle_configuration.request.lifecycle: Still creating... [01m00s elapsed]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Still creating... [01m00s elapsed]
aws_s3_bucket.lifecycle_configuration.request.lifecycle: Still creating... [01m00s elapsed]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Creation complete after 1m11s [id=eugene-unique-request-bucket]
aws_s3_bucket.lifecycle_configuration.response.lifecycle: Creation complete after 1m19s [id=eugene-unique-response-bucket]

Apply complete! Resources: 10 added, 0 changed, 0 destroyed.
```

The screenshot confirms successful creation of the required S3 buckets during Terraform deployment.

Buckets are containers for data stored in S3.

Name	AWS Region	Creation date
<a href="#">cf-templates-qnj7ls6t84aj-us-east-1</a>	US East (N. Virginia) us-east-1	July 23, 2025, 16:04:04 (UTC+00:00)
<a href="#">eugene-unique-request-bucket</a>	US East (N. Virginia) us-east-1	September 3, 2025, 13:07:13 (UTC+00:00)
<a href="#">eugene-unique-response-bucket</a>	US East (N. Virginia) us-east-1	September 3, 2025, 13:07:13 (UTC+00:00)

The eugene-unique-request-bucket has been successfully created. No objects have been uploaded yet.

**eugene-unique-request-bucket** [Info](#)

**Objects (0)**

Name	Type	Last modified	Size	Storage class
------	------	---------------	------	---------------

No objects  
You don't have any objects in this bucket.

[Upload](#)

The eugene-unique-response-bucket has been successfully created and is ready for use

The screenshot shows the AWS S3 console for the bucket 'eugene-unique-response-bucket'. The top navigation bar includes tabs for Objects, Metadata, Properties, Permissions, Metrics, Management, and Access Points. The 'Objects' tab is selected. Below the tabs, there's a toolbar with actions like Copy S3 URI, Copy URL, Download, Open, Delete, Actions, and Create folder. A prominent orange 'Upload' button is visible. A search bar allows filtering by prefix. A table header row shows columns for Name, Type, Last modified, Size, and Storage class. The main content area displays a message stating 'No objects' and 'You don't have any objects in this bucket.'

The translate-json-files Lambda function has been successfully deployed using a ZIP package with Python 3.9 runtime.

The screenshot shows the AWS Lambda console. The top navigation bar includes tabs for Functions, Triggers, and Policies. The 'Functions' tab is selected. It displays a list of functions with one item: 'translate-json-files'. The function details show it was last fetched 4 minutes ago, has a status of 'Active', and was created 7 minutes ago. The function name is 'translate-json-files', the package type is 'Zip', and the runtime is 'Python 3.9'. A large orange 'Create function' button is visible. A search bar at the top allows filtering by attributes or keyword. A table header row shows columns for Function name, Description, Package type, Runtime, and Last modified. The main content area displays the function details.

The translate-json-files Lambda function is successfully configured with an S3 bucket trigger.

Lambda > Functions > translate-json-files

## translate-json-files

**Function overview** [Info](#)

[Throttle](#) [Copy ARN](#) [Actions ▾](#)

[Export to Infrastructure Composer](#) [Download ▾](#)

**Diagram** | **Template**

**translate-json-files**

**S3**

[+ Add destination](#)

[+ Add trigger](#)

**Description**  
-

**Last modified**  
8 minutes ago

**Function ARN**  
arn:aws:lambda:us-east-1:611837360746:function:translate-json-files

**Function URL** | [Info](#)  
-

Confirmed upload of sample\_request.json to the eugene-unique-request-bucket via AWS CLI.

```
PS C:\Users\quaye\Desktop\Capstone Project 2> aws s3 cp "C:\Users\quaye\Desktop\Capstone Project 2\requests" s3://eugene-unique-request-bucket/requests/ --recursive
>>
upload: requests\sample_request.json to s3://eugene-unique-request-bucket/requests/sample_request.json
PS C:\Users\quaye\Desktop\Capstone Project 2>
```

The requests/ folder has been successfully created in the eugene-unique-request-bucket.

eugene-unique-request-bucket [Info](#)

[Objects](#) [Metadata](#) [Properties](#) [Permissions](#) [Metrics](#) [Management](#) [Access Points](#)

**Objects (1)**

[Upload](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

[Name](#) [▲](#) [Type](#) [▼](#) [Last modified](#) [▼](#) [Size](#) [▼](#) [Storage class](#) [▼](#)

Name	Type	Last modified	Size	Storage class
<a href="#">requests/</a>	Folder	-	-	-

The sample\_request.json file is successfully stored in the requests/ folder of the eugene-unique-request-bucket.

The screenshot shows the AWS S3 console with the bucket 'eugene-unique-request-bucket'. The 'Objects' tab is selected. In the 'Objects (1)' section, there is one item: 'sample\_request.json' (Type: json). The object was last modified on September 3, 2025, at 14:23:41 (UTC+00:00) and has a size of 295.0 B. The storage class is Standard. There are buttons for Copy S3 URI, Copy URL, Download, Open, Delete, Actions, and Create folder. A search bar at the top says 'Find objects by prefix' and shows the number '1'. The table headers are Name, Type, Last modified, Size, and Storage class.

Name	Type	Last modified	Size	Storage class
sample_request.json	json	September 3, 2025, 14:23:41 (UTC+00:00)	295.0 B	Standard

The responses/ folder has been successfully created in the eugene-unique-response-bucket.

The screenshot shows the AWS S3 console with the bucket 'eugene-unique-response-bucket'. The 'Objects' tab is selected. In the 'Objects (1)' section, there is one item: 'responses/' (Type: Folder). The object was last modified on '-'. There are buttons for Copy S3 URI, Copy URL, Download, Open, Delete, Actions, and Create folder. A search bar at the top says 'Find objects by prefix' and shows the number '1'. The table headers are Name, Type, Last modified, Size, and Storage class.

Name	Type	Last modified	Size	Storage class
responses/	Folder	-	-	-

The translated file sample\_request\_translated.json has been successfully stored in the responses/ folder of the eugene-unique-response-bucket.

The screenshot shows the AWS S3 console interface. At the top, there's a breadcrumb navigation 'responses/'. On the right, there's a blue button labeled 'Copy S3 URI'. Below this, there are two tabs: 'Objects' (which is selected) and 'Properties'. Under the 'Objects' tab, there's a heading 'Objects (1)'. Below it is a toolbar with several buttons: 'Copy S3 URI', 'Copy URL', 'Download', 'Open', 'Delete', 'Actions', and 'Create folder'. A prominent orange 'Upload' button is located below the toolbar. A search bar with placeholder text 'Find objects by prefix' is present. To the right of the search bar are navigation arrows and a gear icon. The main table lists the single object: 'sample\_request\_translated.json' (Type: json, Last modified: September 3, 2025, Size: 545.0 B, Storage class: Standard).

Name	Type	Last modified	Size	Storage class
sample_request_translated.json	json	September 3, 2025, 14:23:44 (UTC+00:00)	545.0 B	Standard

An auto-delete lifecycle rule has been successfully enabled for the eugene-unique-request-bucket.

The screenshot shows the AWS S3 console interface under the 'Management' tab for the 'eugene-unique-request-bucket'. In the 'Lifecycle configuration' section, there's a note about managing objects cost effectively and a link to learn more. It also shows the minimum object size for transitions (All storage classes 128K). In the 'Lifecycle rules (1)' section, there's a table with one rule named 'auto-delete'. The table columns include Lifecycle rule name, Status, Scope, Current version actions, Noncurrent versions actions, Expired object delete marker actions, and Incomplete multipart upload actions. The 'auto-delete' rule is enabled and applies to the entire bucket, with its current action being 'Expires'. There's also a link to 'View lifecycle configuration'.

Lifecycle rule name	Status	Scope	Current version actions	Noncurrent versions actions	Expired object delete marker actions	Incomplete multipart upload actions
auto-delete	Enabled	Entire bucket	Expires	-	-	-

The auto-delete lifecycle rule is configured to expire objects in the bucket 7 days after upload.

The screenshot shows the 'Lifecycle rule configuration' for the 'auto-delete' rule. It includes sections for 'Prefix', 'Object tags', 'Minimum object size' (with a note about determining size for transitions), and 'Maximum object size'. Below this, a 'Review transition and expiration actions' section details actions at Day 0 (Objects uploaded) and Day 7 (Objects expire). A downward arrow indicates the progression from Day 0 to Day 7.

Day 0	Day 7
• Objects uploaded	• Objects expire

An auto-delete lifecycle rule is active for the eugene-unique-response-bucket, set to expire current version objects.

The screenshot shows the 'Management' tab for the 'eugene-unique-response-bucket'. Under 'Lifecycle configuration', it explains what lifecycle rules do and defines a 'Default minimum object size for transitions' of 128K. The 'Lifecycle rules (1)' section shows one rule named 'auto-delete' which is enabled and applies to the entire bucket, with an 'Expires' action set to 7 days. A link to 'View lifecycle configuration' is also present.

Lifecycle rule name	Status	Scope	Current version actions	Noncurrent versions actions	Expired object delete ma...	Incomplete multipart up...
auto-delete	Enabled	Entire bucket	Expires	-	-	-

The auto-delete lifecycle rule for the response bucket is set to expire objects 7 days after upload.

The screenshot shows the 'auto-delete' lifecycle rule configuration for a bucket. At the top right are 'Edit', 'Delete', and 'Actions' buttons. The configuration section includes fields for 'Lifecycle rule name' (auto-delete), 'Status' (Enabled), and 'Scope' (Entire bucket). It also lists 'Prefix', 'Object tags', and sections for 'Minimum object size' and 'Maximum object size'. Below this is a 'Review transition and expiration actions' section with 'Current version actions' and 'Noncurrent versions actions' tables. The 'Current version actions' table shows actions at 'Day 0' (Objects uploaded) and 'Day 7' (Objects expire). The 'Noncurrent versions actions' table shows 'Day 0' with 'No actions defined.'

An event notification has been set up to trigger the translate-json-files Lambda function on all object creation events in the requests/ folder.

The screenshot shows the 'Event notifications (1)' configuration. At the top right are 'Edit', 'Delete', and a 'Create event notification' button. The table lists one notification rule: 'tf-s3-lambda-20250903130731761800000001' which triggers 'All object create events' with 'requests/, json' filters, 'Lambda function' destination, and 'translate-json-files' as the target.

The lambda\_translate\_exec role has been actively used by the Lambda function, confirming successful role assignment.

Roles (6) <a href="#">Info</a>		
<input type="checkbox"/>	Role name	▲ Trusted entities   Last activity
<input type="checkbox"/>	<a href="#">AWSServiceRoleForSupport</a>	AWS Service: support (Service-Linked) -
<input type="checkbox"/>	<a href="#">AWSServiceRoleForTrustedAdvisor</a>	AWS Service: trustedadvisor (Service) -
<input type="checkbox"/>	<a href="#">first_role_created</a>	Account: 611837360746 -
<input type="checkbox"/>	<a href="#">lambda_translate_exec</a>	AWS Service: lambda 1 hour ago
<input type="checkbox"/>	<a href="#">rds-monitoring-role</a>	AWS Service: monitoring.rds -
<input type="checkbox"/>	<a href="#">second_role_created</a>	Account: 611837360746 102 days ago

The lambda\_translate\_policy grants necessary permissions for Lambda to access Amazon Translate and interact with designated S3 buckets.

<a href="#">lambda_translate_policy</a> <a href="#">Info</a>		<a href="#">Edit</a>	<a href="#">Delete</a>																				
Allows Lambda to use Translate and access specific S3 buckets																							
Policy details																							
Type Customer managed	Creation time September 03, 2025, 13:07 (UTC)	Edited time September 03, 2025, 13:07 (UTC)	ARN <a href="#">arn:aws:iam::611837360746:policy/lambda_translate_policy</a>																				
<a href="#">Permissions</a>	<a href="#">Entities attached</a>	<a href="#">Tags</a>	<a href="#">Policy versions (1)</a>																				
			<a href="#">Last Accessed</a>																				
Permissions defined in this policy <a href="#">Info</a>																							
Permissions defined in this policy document specify which actions are allowed or denied. To define permissions for an IAM identity (user, user group, or role), attach a policy to it																							
<table border="1"> <thead> <tr> <th colspan="4">Allow (3 of 450 services)</th> </tr> <tr> <th>Service</th> <th>▲ Access level</th> <th>▼ Resource</th> <th>Request condition</th> </tr> </thead> <tbody> <tr> <td>CloudWatch Logs</td> <td>Limited: Write</td> <td>region  string like  All</td> <td>None</td> </tr> <tr> <td>S3</td> <td>Limited: List, Read, Write</td> <td>Multiple</td> <td>None</td> </tr> <tr> <td>Translate</td> <td>Limited: Read</td> <td>All resources</td> <td>None</td> </tr> </tbody> </table>				Allow (3 of 450 services)				Service	▲ Access level	▼ Resource	Request condition	CloudWatch Logs	Limited: Write	region  string like  All	None	S3	Limited: List, Read, Write	Multiple	None	Translate	Limited: Read	All resources	None
Allow (3 of 450 services)																							
Service	▲ Access level	▼ Resource	Request condition																				
CloudWatch Logs	Limited: Write	region  string like  All	None																				
S3	Limited: List, Read, Write	Multiple	None																				
Translate	Limited: Read	All resources	None																				
<a href="#">Edit</a> <a href="#">Summary</a> <a href="#">JSON</a>																							
<span>Show remaining 447 services</span>																							

## Testing & Validation

Sample Input (requests/sample\_request.json):

```
requests > {} sample_request.json > ...
1  {
2    "source_lang": "en",
3    "target_lang": "es",
4    "texts": [
5      "Hello, world! It's a beautiful day to explore new ideas and build something amazing!",
6      "This capstone project is truly impressive. It's challenging, exciting, and a great way to showcase everything we've learned"
7    ]
8 }
```

Expected Output (responses/sample\_request\_translated.json):

```
C > Users > quaye > Downloads > {} sample_request_translated.json > ...
1 ✓ {"input_language": "en", "output_language": "es",
2   "original":
3     ["Hello, world! It's a beautiful day to explore new ideas and build something amazing!",
4      "This capstone project is truly impressive. It's challenging, exciting, and a great way to showcase everything we've learned"],
5
6   "translated": ["\u00a1Hola, mundo! \u00a1Es un hermoso d\u00f3nada para explorar nuevas ideas y construir algo incre\u00e1edible!",
7      "\u201cEste proyecto final es realmente impresionante. Es desafiante, emocionante y una excelente manera de mostrar todo lo que hemos aprendido"]
8 }
```

- Use CLI to upload input files.
- Check Lambda logs in CloudWatch for errors.
- Download response files for validation.

The terraform destroy command was executed, beginning the teardown of all previously provisioned AWS resources.

```
PS C:\Users\quaye\Desktop\Capstone Project 2> terraform destroy
data.aws_iam_policy_document.lambda_assume_role: Reading...
aws_s3_bucket.request: Refreshing state... [id=eugene-unique-request-bucket]
aws_s3_bucket.response: Refreshing state... [id=eugene-unique-response-bucket]
data.aws_iam_policy_document.lambda_assume_role: Read complete after 0s [id=2690255455]
aws_iam_role.lambda_exec: Refreshing state... [id=lambda_translate_exec]
aws_lambda_function.translate_function: Refreshing state... [id=translate-json-files]
aws_lambda_permission.allow_s3: Refreshing state... [id=AllowExecutionFromS3]
aws_s3_bucket_notification.request_trigger: Refreshing state... [id=eugene-unique-request-bucket]
aws_s3_bucket_lifecycle_configuration.request.lifecycle: Refreshing state... [id=eugene-unique-request-bucket]
aws_iam_policy.lambda_policy: Refreshing state... [id=arn:aws:iam::611837360746:policy/lambda_translate_policy]
aws_s3_bucket_lifecycle_configuration.response.lifecycle: Refreshing state... [id=eugene-unique-response-bucket]
```

All provisioned AWS resources, including S3 buckets, Lambda function, and IAM roles, have been successfully destroyed.

```
aws_s3_bucket.request: Destroying... [id=eugene-unique-request-bucket]
aws_s3_bucket.response: Destroying... [id=eugene-unique-response-bucket]
aws_lambda_function.translate_function: Destruction complete after 0s
aws_iam_role.lambda_exec: Destroying... [id=lambda_translate_exec]
aws_iam_role.lambda_exec: Destruction complete after 1s
aws_s3_bucket.request: Destruction complete after 2s
aws_s3_bucket.response: Destruction complete after 3s
```

**Destroy complete! Resources: 10 destroyed.**

## Challenges

During the development of this serverless AWS NLP Translation Pipeline, several challenges were encountered which provided valuable learning experiences:

- **Lambda Deployment Packaging:** Packaging and uploading the Lambda function as a ZIP file was error-prone, especially regarding path references in Terraform and keeping the deployed code in sync with the latest changes.
- **S3 Bucket Name Collisions:** AWS requires S3 bucket names to be globally unique. During testing, bucket creation sometimes failed due to name collisions, requiring careful planning and dynamic naming strategies.
- **Managing Free Tier Usage:** Close monitoring was needed to avoid exceeding AWS Free Tier quotas, especially during repeated testing, to prevent unexpected charges.

## Lessons learned

This project led to deep insights relevant to both cloud engineering and DevOps best practices:

- **Infrastructure as Code (IaC) is Essential:** Using Terraform brought tremendous value in terms of repeatability, version control, and transparency, enabling fast re-provisioning of environments and straightforward rollback of changes.
- **Security First:** Least privilege IAM policy design is not just best practice, but crucial for real-world deployments. Early investment in security policy design saved significant troubleshooting effort later.
- **Observability is Key:** Effective use of AWS CloudWatch logs proved to be essential for debugging distributed serverless architectures. Proper instrumentation and logging should be included from the start of all cloud projects.
- **Automation Saves Time:** Automating packaging and Lambda deployment processes eliminated many manual errors and ensured continuous delivery of the most up-to-date application code.
- **Resource Naming and State Management:** The globally unique naming requirement for S3 and state management in IaC highlighted the importance of careful resource planning in cloud projects.

## Cost management

Throughout the project, careful attention was paid to control costs and remain within the AWS Free Tier:

- **Lifecycle Policies:** Implemented on all S3 buckets to automatically expire and delete data after seven days, minimizing storage costs and preventing orphaned resources.
- **Free Tier Resource Usage:** Selected only services included in AWS Free Tier (Lambda, S3, Translate) and kept overall usage well below monthly limits by using small sample files and limiting translation operations.
- **Minimal Permissions:** Used narrowly scoped IAM policies and roles, reducing the risk of accidental access to costly AWS services outside of the project's requirements.
- **Regular Teardown:** Enforced a practice of destroying infrastructure (terraform destroy) immediately upon project completion or after test iterations, ensuring no lingering or idle resources run up costs.
- **Monitoring:** Utilized the AWS Billing Dashboard for continuous oversight of resource consumption and early detection of any potential overages.

## Conclusion

This capstone project demonstrates the successful design and deployment of a secure, cost-effective, and scalable serverless NLP translation pipeline using AWS and Terraform. Leveraging infrastructure as code and managed cloud services enables extremely rapid iteration, high availability, and robust security without incurring the overhead of traditional infrastructure management.

The pipeline automates the process from resource provisioning, through translation workflows, to secure data storage, with full traceability and cost control. The experiences and skills gained through this project—including proficiency in Infrastructure as Code (IaC), serverless computing, cloud security, and cost-effective cloud engineering—provide a strong foundation for developing modern cloud-native solutions and contribute valuable preparation for future real-world deployments.

## References

- AWS Translate: <https://docs.aws.amazon.com/translate/latest/dg/what-is.html>
- Terraform AWS Provider: <https://registry.terraform.io/providers/hashicorp/aws/latest/docs>
- AWS Lambda Python: <https://docs.aws.amazon.com/lambda/latest/dg/python-handler.html>
- AWS S3: <https://docs.aws.amazon.com/s3/>